



J.A.R.V.I.S

An AI assistant that reasons like a **team** — and runs on your own machine.

Jarvis takes a request by voice or text and carries it all the way to a finished action on your behalf — replying to an email, booking a meeting, or pulling a market briefing. Instead of one model guessing, a pipeline of specialised agents routes the request, plans the steps, checks its own work and scores the result before you ever see it. It runs a private model locally at no cost, and switches to a powerful cloud model with a single setting.



Understands & acts

Speak or type. Jarvis transcribes, reads intent, and performs the real action — not just a chat reply.



Reasons in stages

Five specialised agents plan, execute, critique and score every response before it reaches you.



Private by default

Runs a local model at \$0 with data kept on-device. One setting flips the whole stack to the cloud.



Connected to your world

Gmail, Calendar, markets, news, files, music and more — a dozen integrations the agents call on demand.

AT A GLANCE

Local-first · **\$0** to run

5-agent reasoning pipeline


12+ live integrations

Voice + text, hands-free

Local or cloud, one setting

Six layers, one continuous path

The full stack, top to bottom. A request enters at the top and flows down through each layer until it becomes a real action.



THE THROUGH-LINE

Each layer hands its result to the next. Memory is read when the request is routed and written after every turn, so context — who you mean and what just happened — carries across the whole conversation.

How a single request **travels**

A worked example, traced through every layer, so the hand-offs and connections are explicit.

YOU SAY *“Reply to email 1 with a thumbs-up, and remind me to follow up at 5.”*

01 Capture

Silero VAD detects the instant you stop talking, so you're never cut off mid-sentence.

02 Transcribe

Whisper turns speech into text on-device. The audio never leaves your machine.

03 Route

The Router reads intent and spots two actions, tagging them as fast, tool-only tasks.

04 Plan

The Planner splits the request into ordered steps: reply to email #1, then set a 5pm reminder.

05 Act

The Gmail agent sends a threaded thumbs-up reply; the reminder store schedules the 5pm nudge.

06 Verify

The Critic and Evaluator confirm both steps succeeded and score the result.

07 Compose

The Summariser writes one short, natural confirmation of what was done.

08 Respond

ElevenLabs speaks it back while the dashboard updates your inbox and reminders.

Throughout, Memory supplies the context that makes this possible — resolving who ‘email 1’ is, and recording the outcome so the next request is smarter.

A team of specialised agents

Rather than one model improvising, each agent owns one job. They run in sequence, and the chain can loop back to re-plan if quality is low.



Router
CLASSIFIES THE REQUEST

Sends simple tasks straight to a tool and hard ones to the full pipeline — saving time and tokens.



Planner
BREAKS WORK INTO STEPS

Turns a vague ask into concrete, executable steps (a DAG) with the right tools attached to each.



Critic
SANITY-CHECKS THE PLAN

Catches weak or risky plans before any action runs, and can trigger a re-plan.



Evaluator
SCORES THE RESULT

Grades every answer 0–1 so low-quality results are caught, logged and improved before shipping.



Summariser
WRITES THE FINAL REPLY

Produces the single, natural response you actually see or hear, with nothing extra.



Memory
REMEMBERS ACROSS TURNS

A ChromaDB vector store that carries context between requests, so Jarvis isn't starting cold each time.

THE ENGINE BEHIND THEM — ONE VARIABLE SWAPS THE WHOLE STACK

 **Ollama · Local**

DEFAULT

Llama 3.2 on your own hardware. Private, offline-capable and free to run.

 **Groq · Cloud**

ONE SETTING

Llama 3.3 70B via Groq, deployed on Fly.io. More power for heavier reasoning when you want it.

What it can actually do

Twelve live integrations the agents call to take action, plus a dedicated financial analyst, and how it deploys.



Gmail

Read, search and reply in-thread; sends real email for you.



Calendar

Create and read Google Calendar events; checks availability.



Web Search

Pulls live answers from the web for current questions.



News

Aggregates headlines from 25+ sources into briefings.



Markets

Live stock and crypto prices for NVDA, AAPL, BTC and more.



Sports

Scores and tables across football, basketball and cricket.



Weather

Current conditions and forecasts for any location.



Spotify

Play, pause, skip and control playback by voice.



Mac Control

Open apps, set volume and brightness, take screenshots.



File Manager

Find, read and organise files across Desktop, Docs, Downloads.



Reminders

Natural-language timed reminders, like 'in 30 minutes'.



Prayer Times

Accurate daily prayer times for your location.



FinEx — the financial-statement sub-agent

A specialist that ingests company annual reports (PDF) into a Postgres database, then answers financial questions over them by writing SQL and reasoning with an LLM — returning grounded, cited answers rather than guesses.

Deployment & privacy. Runs locally on macOS for full privacy at zero cost, or deploys to the cloud (Fly.io + Groq + Neon Postgres) for a shareable URL. Mac-only features degrade gracefully in the cloud.

Stack: FastAPI · asyncio · ChromaDB · Whisper · Silero VAD · ElevenLabs · Ollama · Groq · Postgres